

# 2 - Search Algorithms

---

## 知识点 & 题目

---

### Basic state model S(P): Classical Planning

- finite and discrete state space  $S$
  - a **known initial state**  $s_0 \in S$
  - a set  $S_G \subseteq S$  of goal states
  - actions  $A(s) \subseteq A$  applicable in each  $s \in S$
  - a **deterministic transition function**  $s' = f(a, s)$  for  $a \in A(s)$
  - positive **action costs**  $c(a, s)$
- A solution is a sequence of applicable actions that maps  $s_0$  into  $S_G$ , and it is optimal if it **minimizes sum of action costs** (e.g., # of steps)
  - Different models and controllers obtained by relaxing assumptions in blue

### Blind search vs. heuristic (informed) search:

- Blind search algorithms: Only use the basic ingredients for general search algorithms.
  - e.g., Depth First Search (DFS), Breadth-first search (BrFS), Uniform Cost (Dijkstra), Iterative Deepening (ID)
- Heuristic search algorithms: Additionally use heuristic functions which estimate the distance (or remaining cost) to the goal.
  - e.g.,  $A^*$ , IDA, Hill Climbing, Best First, WA, DFS B&B, LRTA, . . .
- For satisficing planning, heuristic search vastly outperforms blind algorithms pretty much everywhere.
- For optimal planning, heuristic search also is better (but the difference is less pronounced).
- Blind search does not require any input beyond the problem.
  - Pro: No additional work for the programmer.
  - Con: It's not called "blind" for nothing . . . same expansion order regardless what the problem actually is. Rarely effective in practice.
- Informed search requires as additional input a heuristic function  $h$  that maps states to estimates of their goal distance.
  - Pro: Typically more effective in practice.
  - Con: Implement  $h$ .
  - Note: In planning,  $h$  is generated automatically from the declarative problem description.

## Systematic search vs. local search

- Systematic search algorithms: Consider a large number of search nodes simultaneously.
- Local search algorithms: Work with one (or a few) candidate solutions (search nodes) at a time.
- This is not a black-and-white distinction; there are crossbreeds (e.g., enforced hill-climbing).
- For satisficing planning, there are successful instances of each.
- For optimal planning, systematic algorithms are required.

## Search Terminology

**Search node  $n$ :** Contains a *state* reached by the search, plus information about how it was reached.

**Path cost  $g(n)$ :** The cost of the path reaching  $n$ .

**Optimal cost  $g^*$ :** The cost of an optimal solution path. For a state  $s$ ,  $g^*(s)$  is the cost of a cheapest path reaching  $s$ .

**Node expansion:** Generating all successors of a node, by applying all actions applicable to the node's state  $s$ . Afterwards, the *state*  $s$  itself is also said to be expanded.

**Search strategy:** Method for deciding which node is expanded next.

**Open list:** Set of all *nodes* that currently are candidates for expansion. Also called **frontier**.

**Closed list:** Set of all *states* that were already expanded. Used only in **graph search**, not in **tree search** (up next). Also called **explored set**.

## Search States vs. Search Nodes

- **Search states  $s$ :** States (vertices) of the search space.
- **Search nodes  $\sigma$ :** Search states, plus information on where/when/how they are encountered during search.

### What is in a search node?

Different search algorithms store different information in a search node  $\sigma$ , but typical information includes:

- **$state(\sigma)$ :** Associated search state.
- **$parent(\sigma)$ :** Pointer to search node from which  $\sigma$  is reached.
- **$action(\sigma)$ :** An action leading from  $state(parent(\sigma))$  to  $state(\sigma)$ .
- **$g(\sigma)$ :** Cost of  $\sigma$  (cost of path from the root node to  $\sigma$ ).

For the root node,  $parent(\sigma)$  and  $action(\sigma)$  are undefined.

## Criteria for evaluating search strategies

### Guarantees:

**Completeness:** Is the strategy guaranteed to find a solution when there is one?

**Optimality:** Are the returned solutions guaranteed to be optimal?

### Complexity:

**Time Complexity:** How long does it take to find a solution? (Measured in **generated states**.)

**Space Complexity:** How much memory does the search require? (Measured in **states**.)

### Typical state space features governing complexity:

**Branching factor  $b$ :** How many successors does each state have?

**Goal depth  $d$ :** The number of actions required to reach the shallowest goal state.

## Blind Search

- Breadth-First Search
- Depth-First Search
- Iterative Deepening Search

## Heuristic Search: Systematic

- Greedy best-first search
- A\*
- Weighted A\*
- Iterative deepening A\* (IDA\*)
- Bidirectional A\* Enhanced (BAE\*)

## Heuristic Search: Local

- Hill-climbing
- Enforced hill-climbing
- Beam search, tabu search, genetic algorithms, simulated annealing, . . .

## Heuristic Function L2 P25 T3

- Heuristic function  $h$  estimates the cost of an optimal path to the goal.
  - Search gives a preference to explore states with small  $h$ .
- Remaining cost:
  - The perfect heuristic  $h^*$ , assigns every state its remaining cost as the heuristic value.
- Search performance depends crucially on the **informedness of  $h$**  and on the **computational overhead** of computing  $h$ .
- Extreme cases:

- $h = h^*$ : Perfectly informed; computing it = solving the planning task in the first place.
  - $h = 0$ : No information at all; can be “computed” in constant time.
- Successful heuristic search requires a good trade-off between  $h$ ’s informedness and the computational overhead of computing it.
- Devise methods that yield good estimates at reasonable computational costs.
- **Definition (Safe/Goal-Aware/Admissible/Consistent).** Let  $\Pi$  be a planning task with state space  $\Theta_{\Pi} = (S, L, c, T, I, S^G)$ , and let  $h$  be a heuristic for  $\Pi$ . The heuristic is called:
  - **safe** if  $h^*(s) = \infty$  for all  $s \in S$  with  $h(s) = \infty$ ;
  - **goal-aware** if  $h(s) = 0$  for all goal states  $s \in S^G$ ;
  - **admissible** if  $h(s) \leq h^*(s)$  for all  $s \in S$ ;
  - **consistent** if  $h(s) \leq h(s') + c(a)$  for all transitions  $s \xrightarrow{a} s'$ .
- If  $h$  is consistent and goal-aware, then  $h$  is admissible.
- If  $h$  is admissible, then  $h$  is goal-aware.
- If  $h$  is admissible, then  $h$  is safe.
- No other implications of this form hold.

## Breadth-First Search L2 P14 T1

- Strategy: Expand nodes in the order they were produced (FIFO frontier).
- Completeness: yes
- Optimality:
  - Yes, for uniform action costs. Breadth-first search always finds a shallowest goal state.
  - If costs are not uniform, this is not necessarily optimal.
- Complexity:

**Time Complexity:** Say that  $b$  is the maximal branching factor, and  $d$  is the goal depth (depth of shallowest goal state).

- **Upper bound on the number of generated nodes?**  $b + b^2 + b^3 + \dots + b^d$ : In the worst case, the algorithm generates all nodes in the first  $d$  layers.
- So the time complexity is  $O(b^d)$ .
- **And if we were to apply the goal test at node-expansion time, rather than node-generation time?**  $O(b^{d+1})$  because then we’d generate the first  $d + 1$  layers in the worst case.

**Space Complexity:** Same as time complexity since all generated nodes are kept in memory.

- Which is the worse problem, time or memory?
  - Typically exhausts RAM memory within a few minutes.
- Breadth-first search is optimal but uses exponential space.

## Depth-First Search L2 P17 T1

- Strategy: Expand the most recent nodes in (LIFO frontier).
- Completeness:
  - No, because search branches may be infinitely long: No check for cycles along a branch!
  - Depth-first search is complete in case the state space is acyclic, e.g., Constraint Satisfaction Problems. If we do add a cycle check, it becomes complete for finite state spaces.
- Optimality: No. After all, the algorithm just “chooses some direction and hopes for the best”.
  - Depth-first search is a way of “hoping to get lucky”.
- Complexity:
  - Space: Stores nodes and applicable actions on the path to the current node. So if  $m$  is the maximal depth reached, the complexity is  **$O(bm)$** .
  - Time: If there are paths of length  $m$  in the state space,  **$O(b^m)$**  nodes can be generated. Even if there are solutions of depth 1!
    - If we happen to choose “the right direction” then we can find a length-1 solution in time  $O(b)$  regardless how big the state space is.
- Depth-first search uses linear space but is not optimal.

## Iterative Deepening Search L2 P19 T1

- Completeness: yes
- Optimality: yes
- Complexity:
  - Space:  $O(bd)$
  - Time:

|                            |   |
|----------------------------|---|
| Breadth-First-Search       | $b + b^2 + \dots + b^{d-1} + b^d \in O(b^d)$                      |
| Iterative Deepening Search | $(d)b + (d-1)b^2 + \dots + 3b^{d-2} + 2b^{d-1} + 1b^d \in O(b^d)$ |

- IDS combines the advantages of breadth-first and depth-first search.
- It is the preferred blind search method in large state spaces with unknown solution depth.

## Greedy best-first search T2



## Greedy Best-First Search (with duplicate detection)

```
open := new priority queue ordered by ascending  $h(\text{state}(\sigma))$ 
open.insert(make-root-node(init()))
closed :=  $\emptyset$ 
while not open.empty():
     $\sigma := \text{open.pop-min}()$  /* get best state */
    if  $\text{state}(\sigma) \notin \text{closed}$ : /* check duplicates */
        closed := closed  $\cup \{\text{state}(\sigma)\}$  /* close state */
        if is-goal(state( $\sigma$ )): return extract-solution( $\sigma$ )
        for each  $(a, s') \in \text{succ}(\text{state}(\sigma))$ : /* expand state */
             $\sigma' := \text{make-node}(\sigma, a, s')$ 
            if  $h(\text{state}(\sigma')) < \infty$ : open.insert( $\sigma'$ )
return unsolvable
```

- Completeness: Yes, for safe heuristics. (and duplicate detection to avoid cycles)
- Optimality: No
- Invariant under all strictly monotonic transformations of  $h$ 
  - e.g., scaling with a positive constant or adding a constant.
- Priority queue: e.g., a min heap.
- “Check Duplicates”: Could already do in “expand state”; done here after “get best state” only to more clearly point out relation to A.

## A\* L2 P33 T2

### A\* (with duplicate detection and re-opening)

```
open := new priority queue ordered by ascending  $g(\text{state}(\sigma)) + h(\text{state}(\sigma))$ 
open.insert(make-root-node(init()))
closed :=  $\emptyset$ 
best-g :=  $\emptyset$  /* maps states to numbers */
while not open.empty():
     $\sigma := \text{open.pop-min}()$ 
    if  $\text{state}(\sigma) \notin \text{closed}$  or  $g(\sigma) < \text{best-g}(\text{state}(\sigma))$ :
        /* re-open if better  $g$ ; note that all  $\sigma'$  with same state but worse  $g$ 
           are behind  $\sigma$  in open, and will be skipped when their turn comes */
        closed := closed  $\cup \{\text{state}(\sigma)\}$ 
        best-g(state( $\sigma$ )) :=  $g(\sigma)$ 
        if is-goal(state( $\sigma$ )): return extract-solution( $\sigma$ )
        for each  $(a, s') \in \text{succ}(\text{state}(\sigma))$ :
             $\sigma' := \text{make-node}(\sigma, a, s')$ 
            if  $h(\text{state}(\sigma')) < \infty$ : open.insert( $\sigma'$ )
return unsolvable
```

- $f$ -value of a state: defined by  $f(s) := g(s) + h(s)$ .
- Generated nodes: Nodes inserted into open at some point.
- Expanded nodes: Nodes popped from open for which the test against closed and distance succeeds.

- Re-expanded nodes: Expanded nodes for which state is closed upon expansion (also called re-opened nodes).
- Completeness: Yes, for safe heuristics. (Even without duplicate detection.)
- Optimality: Yes, for admissible heuristics. (Even without duplicate detection.)
- Popular method: break ties ( $f(s) = f(s')$ ) by smaller  $h$ -value.
- If  $h$  is admissible and consistent, then A never re-opens a state. So if we know that this is the case, then we can simplify the algorithm.
- Common, hard to spot bug: check duplicates at the wrong point.

## Weighted A\* T2

### Weighted A\* (with duplicate detection and re-opening)

```

open := new priority queue ordered by ascending  $g(\text{state}(\sigma)) + W * h(\text{state}(\sigma))$ 
open.insert(make-root-node(init()))
closed :=  $\emptyset$ 
best-g :=  $\emptyset$ 
while not open.empty():
     $\sigma := \text{open.pop-min}()$ 
    if  $\text{state}(\sigma) \notin \text{closed}$  or  $g(\sigma) < \text{best-g}(\text{state}(\sigma))$ :
        closed := closed  $\cup \{\text{state}(\sigma)\}$ 
        best-g( $\text{state}(\sigma)$ ) :=  $g(\sigma)$ 
        if is-goal( $\text{state}(\sigma)$ ): return extract-solution( $\sigma$ )
        for each  $(a, s') \in \text{succ}(\text{state}(\sigma))$ :
             $\sigma' := \text{make-node}(\sigma, a, s')$ 
            if  $h(\text{state}(\sigma')) < \infty$ : open.insert( $\sigma'$ )
return unsolvable

```

The weight  $W \in \mathbb{R}_0^+$  is an **algorithm parameter**:

- For  $W = 0$ , weighted A\* behaves like uniform-cost search.
- For  $W = 1$ , weighted A\* behaves like A\*.
- For  $W \rightarrow \infty$ , weighted A\* behaves like greedy best-first search.

### Properties:

- For  $W > 1$ , weighted A\* is **bounded suboptimal**: if  $h$  is admissible, then the solutions returned are at most a factor  $W$  more costly than the optimal ones.

## Hill-climbing

## Hill-Climbing

```
 $\sigma := \text{make-root-node}(\text{init}())$ 
forever:
  if  $\text{is-goal}(\text{state}(\sigma))$ :
    return  $\text{extract-solution}(\sigma)$ 
   $\Sigma' := \{ \text{make-node}(\sigma, a, s') \mid (a, s') \in \text{succ}(\text{state}(\sigma)) \}$ 
   $\sigma := \text{an element of } \Sigma' \text{ minimizing } h$  /* (random tie breaking) */
```

### Remarks:

- Makes sense only if  $h(s) > 0$  for  $s \notin S^G$ .
- Is this complete or optimal? No.
- Can easily get stuck in local minima where immediate improvements of  $h(\sigma)$  are not possible.
- Many variations: tie-breaking strategies, restarts, ...

## Enforced hill-climbing A1

### Enforced Hill-Climbing: Procedure *improve*

```
def improve( $\sigma_0$ ):
  queue := new fifo queue
  queue.push-back( $\sigma_0$ )
  closed :=  $\emptyset$ 
  while not queue.empty():
     $\sigma = \text{queue.pop-front}()$ 
    if  $\text{state}(\sigma) \notin \text{closed}$ :
      closed := closed  $\cup \{ \text{state}(\sigma) \}$ 
      if  $h(\text{state}(\sigma)) < h(\text{state}(\sigma_0))$ : return  $\sigma$ 
      for each  $(a, s') \in \text{succ}(\text{state}(\sigma))$ :
         $\sigma' := \text{make-node}(\sigma, a, s')$ 
        queue.push-back( $\sigma'$ )
  fail
```

↪ Breadth-first search for state with strictly smaller  $h$ -value.



## Enforced Hill-Climbing

```
 $\sigma := \text{make-root-node}(\text{init}())$   
while not  $\text{is-goal}(\text{state}(\sigma))$ :  
     $\sigma := \text{improve}(\sigma)$   
return  $\text{extract-solution}(\sigma)$ 
```

### Remarks:

- Makes sense only if  $h(s) > 0$  for  $s \notin S^G$ .
- **Is this optimal?** No.
- **Is this complete?** In general, no. Under particular circumstances, yes. Assume that  $h$  is goal-aware.
  - Procedure *improve* fails: no state with strictly smaller  $h$ -value reachable from  $s$ , thus (with assumption) goal not reachable from  $s$ .
  - This can, for example, not happen if the state space is undirected, i.e., if for all transitions  $s \rightarrow s'$  in  $\Theta_\Pi$  there is a transition  $s' \rightarrow s$ .

## Properties of search algorithms

|          | DFS         | BrFS  | ID          | A*    | HC       | IDA*        |
|----------|-------------|-------|-------------|-------|----------|-------------|
| Complete | No          | Yes   | Yes         | Yes   | No       | Yes         |
| Optimal  | No          | Yes*  | Yes         | Yes   | No       | Yes         |
| Time     | $\infty$    | $b^d$ | $b^d$       | $b^d$ | $\infty$ | $b^d$       |
| Space    | $b \cdot d$ | $b^d$ | $b \cdot d$ | $b^d$ | $b$      | $b \cdot d$ |

- Parameters:  $d$  is solution depth;  $b$  is branching factor
- Breadth First Search (BrFS) optimal when costs are uniform
- A\*/IDA\* optimal when  $h$  is **admissible**;  $h \leq h^*$

## Summary

**Distinguish:** World states, search states, search nodes.

- **World state:** Situation in the world modelled by the planning task.
- **Search state:** Subproblem remaining to be solved.
  - In **progression**, world states and search states are identical.
  - In **regression**, search states are sub-goals describing sets of world states.
- **Search node:** Search state + info on “how we got there”.

**Search algorithms** mainly differ in **order of node expansion**:

- **Blind** vs. **heuristic** (or **informed**) search.
- **Systematic** vs. **local** search.

# 题目

## Quiz

### Question!

If we set  $h(n) := 0$  for all  $n$ , what does  $A^*$  become?

- (A): Breadth-first search. (B): Depth-first search.  
(C): Uniform-cost search. (D): Depth-limited search.

→ (C): Same expansion order. (Details in book-keeping of open/closed states may differ.)

### Question!

If we set  $h(n) := 0$  for all  $n$ , what can greedy best-first search become?

- (A): Breadth-first search. (B): Depth-first search.  
(C): Uniform-cost search. (D): A), B) and C)

→  $h$  implies no ordering of nodes at all, so this fully depends on how we break ties in the open list. (A): FIFO, (B): LIFO, (C): Order on  $g$ . (Details in book-keeping of open/closed states may differ.)

### Question!

Is informed search always better than blind search?

- (A): Yes. (B): No.

→ In greedy best-first search, the heuristic may yield larger search spaces than uniform-cost search. E.g., in path planning, say you want to go from Melbourne to Sydney, but  $h(\text{Perth}) < h(\text{Canberra})$ .

→ In  $A^*$  with an admissible heuristic and duplicate checking, we cannot do worse than uniform-cost search:  $h(s) > 0$  can only reduce the number of states we must consider to prove optimality.

→ Also, in the above example,  $A^*$  doesn't expand Perth with *any* admissible heuristic, because  $g(\text{Perth}) > g(\text{Sydney})$ !

→ "Trusting the heuristic" has its dangers! Sometimes  $g$  helps to reduce search.

I would encourage you to draw a small graph containing 3 nodes: Melbourne, Sydney and Perth. Start the search in Melbourne, and set Sydney as your goal. Draw the costs of travelling from any pair of cities in hours: Melb-Pert = 30h, Perth-Syd=45h, and Melb-Syd=10h. Test the statements above to check your intuitions.

Consider general search problems, which of the following are true?

Correct!

- ☒ Code that implements A\* tree search can be used to run Uniform-cost search.

Yes, you just need to set  $f(n)=g(n)$ .

- ☐ A\* tree search is optimal with any heuristic function.

You Answered

- ☒ A\* tree search is complete with any heuristic function.

Only if the heuristic is safe!

If the heuristic is not safe, then it can assign  $h(n) := \infty$  to all nodes  $n$  that lead to a solution

- ☐ A\* graph search is guaranteed to expand no more nodes than DFS.

Correct!

- ☒ The max of two admissible heuristics is always admissible.

Correct! if  $h_1 \leq h^*$  and  $h_2 \leq h^*$ , then it follows that  $\max(h_1, h_2) \leq h^*$

You Answered

- ☒ A heuristic that always evaluates to  $h(s) = 1$  for non-goal search nodes  $s$  is always admissible.

A simple example where this is not true: If all action costs are 0, then this heuristic is not admissible. Check the video about the properties and relationships that can be proved

**Definition (Safe/Goal-Aware/Admissible/Consistent).** Let  $\Pi$  be a planning task with state space  $\Theta_{\Pi} = (S, L, c, T, I, S^G)$ , and let  $h$  be a heuristic for  $\Pi$ . The heuristic is called:

- **safe** if  $h^*(s) = \infty$  for all  $s \in S$  with  $h(s) = \infty$ ;
- **goal-aware** if  $h(s) = 0$  for all goal states  $s \in S^G$ ;
- **admissible** if  $h(s) \leq h^*(s)$  for all  $s \in S$ ;
- **consistent** if  $h(s) \leq h(s') + c(a)$  for all transitions  $s \xrightarrow{a} s'$ .

→ Relationships?

**Proposition.** Let  $\Pi$  be a planning task with state space  $\Theta_{\Pi} = (S, L, c, T, I, S^G)$ , and let  $h$  be a heuristic for  $\Pi$ . If  $h$  is consistent and goal-aware, then  $h$  is admissible. If  $h$  is admissible, then  $h$  is goal-aware. If  $h$  is admissible, then  $h$  is safe. No other implications of this form hold.



## Question 5

1 / 1 pts

Heuristic functions estimate the distance from a state to:

- ☐ the closest neighbour
- ☐ the initial state
- ☒ the closest goal state
- ☐ the furthest goal state

Correct!

Correct! heuristic functions estimate the cost of solving the problem from the current state, which is the cost to the nearest goal

## Question 6

1 / 1 pts

What is a SAFE heuristic function?

- ☐ If a solution does not exist from state  $s$ , then  $h(s) = \text{infinity}$
- ☒ If a solution exists from state  $s$ , then  $h(s) < \text{infinity}$
- ☐ If a solution does not exist from state  $s$ , the  $h(s) > \text{infinity}$
- ☐ If a solution exists from state  $s$ , then  $h(s) = \text{infinity}$

Correct!

Correct! Safe heuristics ensure that no state will have a heuristic value of infinity if a solution can be found from that state. That way states with a heuristic value of infinity can safely be ignored when searching for a solution.

## Question 7

1 / 1 pts

A goal aware heuristic assigns  $h=0$  to all goal states


- ☒ True
- ☐ False

Correct!

Correct!



Given the graph below, is the heuristic:



☐ Consistent, Goal Aware and Admissible

☐ Goal Aware and Admissible

☐ Goal Aware, Admissible and Safe

Correct!

☒ Consistent, Goal Aware, Admissible and Safe

Consistent: Yes because  $h(s) - h(s') = 0$  for all transitions (as  $h(s)=0$  for all states).  $c(a)=1$  for all actions. Therefore  $h(s)-h(s') = 0 < c(a) = 1$

Goal Aware: Yes because  $h=0$  for state G

Admissible: Yes as  $h(s)=0 \leq h^*(s)$  for all states

Safe: Yes as no states with  $h^*(s) = \text{infinity}$

Heuristic search performance depends on:

Correct!

☒ the informedness of the heuristic and computation overhead

Correct, heuristic performance is always a balance between how well it directs the search (informedness) and how long it takes to compute (computation overhead)

☐ Only on how close our heuristic is to  $h^*$

☐ the day of the month

only on how close our heuristic is to  $h=0$

### Question 1

1 / 1 pts

WA\* uses  $f(n) = g(n) + w \cdot h(n)$ . Then, with  $w=0$  it becomes:

Correct!

- ☐ A\*
- ☒ Uniform-cost search (Dijkstra)
- ☐ Breadth-first search
- ☐ Greedy Best First Search

### Question 2

1 / 1 pts

WA\* with  $w = \text{infinity}$  becomes


Correct!


- ☐ A\*
- ☐ Uniform-cost search (Dijkstra)
- ☐ Breadth-first search
- ☒ Greedy Best First Search

### Question 3

0 / 1 pts

For the following graph, assuming the  $h=0$  node is the goal, Enforced Hill Climbing is guaranteed to find a solution





Correct Answer

☐ Yes

You Answered

☒ No

We can guarantee the search will never get stuck in a state that doesn't lead to the goal, as a path can be found from every state to the goal.