

3 - Introduction to Planning

知识点 & 题目

The key problem is to select the action to do next. This is the so-called control problem. Three approaches to this problem:

- Programming-based: Specify control by hand
- Learning-based: Learn control from experience
- Model-based: Specify problem by hand, derive control automatically

Approaches not orthogonal; successes and limitations in each

Different models yield different types of controllers

Programming-Based Approach

Control specified by programmer

- Advantage: domain-knowledge easy to express
- Disadvantage: cannot deal with situations not anticipated by programmer

Learning-Based Approach

Learns a controller from experience or through simulation

- Unsupervised (Reinforcement Learning):
 - penalize Mario time that 'dies'
 - reward agent each time opponent 'dies' and level is finished
- Supervised (Classification)
 - learn to classify actions into good or bad from info provided by teacher
- Evolutionary
 - from pool of possible controllers: try them out, select the ones that do best, and mutate and recombine for a number of iterations, keeping best
- Advantage: does not require much knowledge in principle
- Disadvantage: in practice, hard to know which features to learn, and is slow

Model-Based Approach / General Problem Solving L3 P10

Specify model for problem: actions, initial situation, goals, and sensors

Let a solver compute controller automatically

Advantage:

- Powerful: In some applications generality is absolutely necessary

- Quick: Rapid prototyping. 10s lines of problem description vs. 1000s lines of C++ code. (Language generation)
- Flexible & Clear: Adapt/maintain the description
- Intelligent & domain-independent: Determines automatically how to solve complex problem effectively!

Disadvantage:

- Efficiency loss: Without any domain-specific knowledge about Chess, you don't beat Kasparov

Trade-off between 'automatic and general' vs. 'manual work but effective'

Model-based approach to intelligent behavior called Planning in AI

Different planning models

- Classical Planning:**
- finite and discrete state space S
 - a **known initial state** $s_0 \in S$
 - a set $S_G \subseteq S$ of goal states
 - actions $A(s) \subseteq A$ applicable in each $s \in S$
 - a **deterministic transition function** $s' = f(a, s)$ for $a \in A(s)$
 - positive **action costs** $c(a, s)$
- A solution is a sequence of applicable actions that maps s_0 into S_G , and it is optimal if it minimizes sum of action costs (e.g., # of steps)
 - Different models and controllers obtained by relaxing assumptions in blue

Conformant Planning:

- finite and discrete state space S
 - a **set of possible initial state** $S_0 \in S$
 - a set $S_G \subseteq S$ of goal states
 - actions $A(s) \subseteq A$ applicable in each $s \in S$
 - a **non-deterministic** transition function $F(a, s) \subseteq S$ for $a \in A(s)$
 - uniform action costs $c(a, s)$
- A solution is still an action sequence but must achieve the goal **for any possible initial state and transition**.
 - More complex than classical planning, verifying that a plan is conformant intractable in the worst case; but special case of planning with partial observability.

Planning with Markov Decision Processes: MDPs are fully observable, probabilistic state models

- a state space S
- initial state $s_0 \in S$
- a set $G \subseteq S$ of goal states
- actions $A(s) \subseteq A$ applicable in each state $s \in S$
- transition probabilities $P_a(s'|s)$ for $s \in S$ and $a \in A(s)$
- action costs $c(a, s) > 0$


- Solutions are functions (policies) mapping states into actions
- Optimal solutions minimize expected cost to goal

Partially Observable MDPs: POMDPs are partially observable, probabilistic state models

- states $s \in S$
- actions $A(s) \subseteq A$
- transition probabilities $P_a(s'|s)$ for $s \in S$ and $a \in A(s)$
- initial belief state b_0
- final belief state b_f
- sensor model given by probabilities $P_a(o|s)$, $o \in Obs$

- Belief states are probability distributions over S
- Solutions are policies that map belief states into actions
- Optimal policies minimize expected cost to go from b_0 to G

Agent **A** must reach **G**, moving one cell at a time in **known** map



- If actions deterministic and initial location known, planning problem is **classical**
- If actions stochastic and location observable, problem is an **MDP**
- If actions stochastic and location partially observable, problem is a **POMDP**

Different combinations of uncertainty and feedback: three problems, three models

Models, Languages, and Solvers

- A planner is a solver over a class of models; it takes a model description, and computes the corresponding controller
- Many models, many solution forms: uncertainty, feedback, costs, . . .
- Models described in suitable planning languages (Strips, PDDL, PPDDL, . . .) where states represent interpretations over the language.

A Basic Language for Classical Planning: STRIPS T4

■ A **problem** in STRIPS is a tuple $P = \langle F, O, I, G \rangle$:

- F stands for set of all **atoms** (boolean vars)
- O stands for set of all **operators** (actions)
- $I \subseteq F$ stands for **initial situation**
- $G \subseteq F$ stands for **goal situation**

■ Operators $o \in O$ **represented** by

- the **Add** list $Add(o) \subseteq F$
- the **Delete** list $Del(o) \subseteq F$
- the **Precondition** list $Pre(o) \subseteq F$


A STRIPS problem $P = \langle F, O, I, G \rangle$ determines **state model** $S(P)$ where

- the states $s \in S$ are **collections of atoms** from F . $S = 2^F$
- the initial state s_0 is I
- the goal states s are such that $G \subseteq s$
- the actions a in $A(s)$ are ops in O s.t. $Pre(a) \subseteq s$
- the next state is $s' = s - Del(a) + Add(a)$
- action costs $c(a, s)$ are all 1

→ (Optimal) **Solution** of P is (optimal) **solution** of $S(P)$

→ Slight language extensions often convenient: **negation**, **conditional effects**, **non-boolean variables**; some required for describing richer models (costs, probabilities, ...).

The Blocks world:



- **Propositions:** $on(x, y)$, $onTable(x)$, $clear(x)$, $holding(x)$, $armEmpty()$.
- **Initial state:** $\{onTable(E), clear(E), \dots, onTable(C), on(D, C), clear(D), armEmpty()\}$.
- **Goal:** $\{on(E, C), on(C, A), on(B, D)\}$.
- **Actions:** $stack(x, y)$, $unstack(x, y)$, $putdown(x)$, $pickup(x)$.
- **$stack(x, y)?$** $pre : \{holding(x), clear(y)\}$
 $add : \{on(x, y), armEmpty(), clear(x)\}$
 $del : \{holding(x), clear(y)\}$.

PDDL T4 A2


PDDL is not a propositional language:

- Representation is lifted, using object variables to be instantiated from a finite set of objects. (Similar to predicate logic)
- Action schemas parameterized by objects.
- Predicates to be instantiated with objects.

A PDDL planning task comes in two pieces:


- The domain file and the problem file.
- The problem file gives the objects, the initial state, and the goal state.
- The domain file gives the predicates and the operators; each benchmark domain has one domain file.

The Blocks World:



- Domain:

```
(define (domain blocksworld)
  (:predicates (clear ?x) (holding ?x) (on ?x ?y)
    (on-table ?x) (arm-empty))
  (:action stack
    :parameters (?x ?y)
    :precondition (and (clear ?y) (holding ?x))
    :effect (and (arm-empty) (on ?x ?y)
      (not (clear ?y)) (not (holding ?x))))
  )
  ...
```



- Problem:

```
(define (problem bw-abcde)
  (:domain blocksworld)
  (:objects a b c d e)
  (:init (on-table a) (clear a)
    (on-table b) (clear b)
    (on-table e) (clear e)
    (on-table c) (on d c) (clear d)
    (arm-empty))
  (:goal (and (on e c) (on c a) (on b d))))
```

Satisficing vs. Optimal

- Satisficing planning
 - Input: A planning task P.
 - Output: A **plan** for P, or 'unsolvable' if no plan for P exists.
 - By **PlanEx**, we denote the problem of deciding, given a planning task P, whether or not there exists a plan for P.
- Optimal planning
 - Input: A planning task P.
 - Output: An **optimal plan** for P, or 'unsolvable' if no plan for P exist.
 - By **PlanLen**, we denote the problem of deciding, given a planning task P and an integer B, whether or not there exists a plan for P of length at most B.
- The techniques successful for either one of these are almost disjoint!
- Satisficing planning is much more effective in practice
- Programs solving these problems are called (optimal) planners, planning systems, or planning tools.

- PlanEx and PlanLen are PSPACE-complete.
 - At least as hard as any other problem contained in PSPACE.
- In general, PlanEx and PlanLen have the same complexity.
- Within particular applications, bounded length plan existence is often harder than plan existence.
- This happens in many IPC benchmark domains: PlanLen is NP-complete while PlanEx is in P.
 - For example: Blocksworld and Logistics.
- In practice, optimal planning is (almost) never 'easy'

NP & PSPACE:

Def Turing machine: Works on a **tape** consisting of **tape cells**, across which its **R/W head** moves. The machine has **internal states**. There are **transition rules** specifying, given the current cell content and internal state, what the subsequent internal state will be, and whether the R/W head moves left or right or remains where it is. Some internal states are **accepting** ('yes'; else 'no').

Def NP: Decision problems for which there exists a non-deterministic Turing machine that runs in time polynomial in the size of its input. Accepts if at least one of the possible runs accepts.

Def PSPACE: Decision problems for which there exists a deterministic Turing machine that runs in space polynomial in the size of its input.

Relation: Non-deterministic polynomial space can be simulated in deterministic polynomial space. Thus **PSPACE = NPSPACE**, and hence (trivially) **NP subset PSPACE**.

Computation

Key issue: exploit two roles of language

- specification: concise model description
- computation: reveal useful heuristic information (structure)

Two traditional approaches: search vs. decomposition

- explicit search of the state model $S(P)$ direct but not effective til recently
- near decomposition of the planning problem thought a better idea

State of the Art in Classical Planning

- significant progress since Graphplan
- empirical methodology
 - standard PDDL language
 - planners and benchmarks available; competitions
 - focus on performance and scalability
- large problems solved (non-optimally)
- different formulations and ideas
 - Planning as Heuristic Search
 - Planning as SAT (Satisfiability)

- Other: Local Search (LPG), Monte-Carlo Search, . . .

Summary

- General problem solving attempts to develop solvers that perform well across a large class of problems.
- Planning, as considered here, is a form of general problem solving dedicated to the class of classical search problems.
 - Actually, we also address inaccessible, stochastic, dynamic, continuous, and multi-agent settings.
- Classical search problems require to find a path of actions leading from an initial state to a goal state.
 - They assume a single-agent, fully-observable, deterministic, static environment. Despite this, they are ubiquitous in practice.
- Heuristic search planning has dominated the International Planning Competition (IPC). We focus on it here.
- STRIPS is the simplest possible, while reasonably expressive, language for our purposes. It uses Boolean variables (facts), and defines actions in terms of precondition, add list, and delete list.
- Plan existence (bounded or not) is PSPACE-complete to decide for STRIPS.
- PDDL is the de-facto standard language for describing planning problems.

题目

Quiz

Question

If planners x, y both compete in IPC'YY, and x wins, is x 'better than' y ?

(A): Yes.

(B): No.

→ Yes, but only on the IPC'YY benchmarks, and only according to the criteria used for determining a 'winner'! On other domains and/or according to other criteria, you may well be better off with the 'looser'.

→ It's complicated, over-simplification is dangerous. (But, of course, nevertheless is being done all the time).

Question 4

1 / 1 pts

Conformant Planning has:

Correct!

- ☐ A set of possible initial states and a probabilistic transition function
- ☒ A set of possible initial states and a non-deterministic transition function
- ☐ A probability distribution over the initial states and a probabilistic transition function
- ☐ A probability distribution over the initial states and a non-deterministic transition function

Question 5

1 / 1 pts

POMDPs have:


Correct!

- ☒ A sensor model given by probabilities drawn from observations about the environment
- ☐ A set of possible initial states
- ☐ A deterministic transition function
- ☐ A complexity that is identical to classical planning

Question 6

0 / 1 pts

Given the following initial state in a Blocks World problem, how many propositions are required to specify the initial state



You Answered

3

Correct Answers

6 (with margin: 0)

The following are required: onTable(A), on(C, A), onTable(B), clear(C), clear(B), armEmpty